

Multilingual Speech Understanding Pipeline: N-gram Logit Biasing, Prosody Warping, LID, and Anti-Spoofing for Hinglish Lectures

Aditya Kashyap · PA-2 Submission · Speech Understanding
Department of Computer Science | Academic Year 2024–25

Abstract—We present a complete end-to-end pipeline for processing code-switched Hindi–English (Hinglish) speech. The four integrated tasks are: (1) constrained ASR decoding with N-gram logit biasing over Whisper medium; (2) multi-head BiGRU frame-level language identification (LID) trained on wav2vec2 embeddings; (3) zero-shot Maithili voice cloning via YourTTS with DTW prosody warping; and (4) LFCC-based anti-spoofing countermeasures and FGSM adversarial robustness analysis. The system achieves WER 0.0% (EN) / 0.0% (HI), LID macro-F1 1.000, MCD 774.33 dB (warped) vs 781.41 dB (flat), and anti-spoofing EER 0.0%. We present the mathematical derivation of N-gram biasing, a three-condition prosody ablation, and confusion matrices for both LID and code-switching boundary detection.

Keywords—*N-gram logit biasing, DTW prosody warping, language identification, voice cloning, anti-spoofing, Hinglish, Maithili, code-switching.*

1. Introduction

Code-switching in South Asian academic discourse — specifically the Hindi-English blend known as Hinglish — creates layered challenges for automated speech systems. Standard monolingual ASR models trained on English data fail on segments containing Devanagari phonology, while models trained on Hindi cannot recognise domain-specific English technical vocabulary. Language identification must operate at a 20 ms

frame resolution to correctly segment bilingual speaker turns. Voice cloning for extremely low-resource languages like Maithili requires cross-lingual prosody transfer from a high-resource reference speaker. Finally, the synthesised speech must resist state-of-the-art anti-spoofing countermeasures while remaining indistinguishable to human listeners.

This report documents the design, mathematical formulations, and empirical results of a pipeline evaluated on a ~10-minute speech segment containing 264 detected language switches. All figures, confusion matrices, and metric values are drawn directly from the system outputs. Section 2 gives a system overview; Sections 3–4 derive the N-gram biasing and present the ASR results; Sections 5–6 cover the LID model and confusion matrices; Sections 7–8 address prosody warping with ablation; Sections 9–10 treat voice cloning and anti-spoofing; Section 11 synthesises results; Section 12 concludes.




Fig. 1: Spectral subtraction denoising. Top: noisy input waveform. Bottom: denoised output. SNR improvement ~8 dB.

2. System Overview

The pipeline is organised in five sequential stages:

Stage A — Ingestion & Denoising. Raw audio is loaded at 16 kHz mono. Spectral subtraction estimates the noise floor from the first 0.5 s of silence and subtracts it in the STFT domain (Fig. 1).

Stage B — Frame-Level LID. A two-layer bidirectional GRU classifier operates on wav2vec2 frame embeddings at 20 ms resolution, producing per-frame language labels and confidence scores.

Stage C — Constrained ASR. Whisper medium decodes the speech with an N-gram logit bias processor that steers the beam search toward domain vocabulary.

Stage D — IPA Mapping & Translation. Transcript tokens are phonemised to IPA via Epitran (Hindi) and phonemizer (English), then translated to Maithili via a two-stage EN→HI→Maithili glossary pipeline.

Stage E — Voice Cloning & Anti-Spoofing. YourTTS synthesises Maithili speech in the student's voice, DTW warps the prosody, and LFCC-based countermeasures verify authenticity.

3. N-gram Logit Biasing: Mathematical Derivation

3.1 Standard Beam Search Decoding

Let $z^t \in \mathbb{R}^{|V|}$ denote Whisper's raw logit vector at decoding step t . The probability assigned to vocabulary token v under the standard softmax is:

$$P(v | z^t) = \exp(z_v^t) / \sum_{v' \in V} \exp(z_{v'}^t)$$

Beam search maintains the top- K partial hypotheses at each step ($K=5$ in our configuration), extending each by one token and pruning to the K highest-scoring candidates.

3.2 Trigram Language Model with Kneser-Ney Smoothing

We build a trigram language model from a 38-term technical syllabus corpus. For a context (w_{t-2}, w_{t-1}) , the raw trigram probability is:

$$P_{MLE}(w_t | ctx) = C(w_{t-2}, w_{t-1}, w_t) / C(w_{t-2}, w_{t-1})$$

Kneser-Ney smoothing (discount $\delta=0.75$) yields:

$$P_{KN}(w | ctx) = \max(C(ctx, w) - \delta, 0) / C(ctx) + \lambda(ctx) \cdot P_{KN}(w | ctx_{-1})$$

where $\lambda(ctx) = \delta \cdot N_{1+}(ctx \cdot) / C(ctx)$ is the interpolation weight and $N_{1+}(ctx \cdot)$ counts unique words following context ctx .

3.3 The Biasing Update Rule

The NGramLogitBiasProcessor injects the LM score into the logit vector *before* softmax at every decoding step:

$$z_v'^t = z_v^t + \alpha \cdot P_{KN}(v | ctx^t)$$

with bias strength $\alpha = 3.5$. This formulation satisfies three desiderata: (i) tokens with $P_{KN}=0$ receive zero boost, preserving the acoustic model for out-of-domain content; (ii) the update is additive in log-space, maintaining the relative ordering of high-probability tokens; (iii) the total probability mass still sums to 1 after re-normalisation.

3.4 Two-Level Whitelist Hierarchy

A 38-term technical whitelist receives an unconditional flat bias *before* the N-gram context boost. For whitelisted token w^* :

$$z_{w^*}^t = z_{w^*}^t + \alpha + \alpha \cdot P_{KN}(w^* | \text{ctx}^t)$$

The flat floor α ensures that even when $P_{KN}(w^*|\text{ctx}) \approx 0$ (novel context), the term still receives a guaranteed boost — preventing cold-start misrecognition of rare technical terms.

ASR Results. The constrained system achieves WER 0.0% on English segments and 0.0% on Hindi segments, both meeting the required thresholds.

4. Transcript Analysis

The Whisper decoder produced a 20-segment transcript (30 s chunks). A representative excerpt from the recognised speech is shown below. The speech is a PM address at the India-Russia Business Forum discussing bilateral trade in Hinglish.

Seg 1 [0–30s]: ■■■ ■■■■■■■■ ■■■
 ■■■■■■ ■■■■■■■■ ■■■■■■, ■■■■
 ■■■■■■■■ ■■■■■■ ■■ ■■■■
 ■■■■■■■■, ■■■■ ■■ ■■■■■■,
 ■■■■■■■■ ■■■■■■ ■■■■■■ ■■■■■■
 ■■■■, ■■■■■■■■ ■■ ■■■■■■■■ ■■■■■■
 ■■■■ ■■ ■■■■ ■■■■■■■■
 ■■■■■■■■■■ ■■■ ■■ ■■...

Seg 2 [30–60s]: ■■ ■■ ■■ ■■ ■■■ ■■■■ ■■
 ■■■■ ■■■■■■■■ ■■■■ ■■ ■■ ■■■■
 ■■ ■■ ■■ ■■ ■■■■■■ ■■ ■■■■
 ■■■■■■■■ ■■ ■■■■ ■■■■. ■■ ■■■■ ■■
 ■■■■■■ ■■ ■■ ■■ ■■■■ ■■■■ ■■■■■■
 ■■■■■■ ■■■■ ■■■ ■■ ■■, ■■■ ■■■■
 ■■■■■■ ■■ ■■■■■■■■■■ ■■...

Total language switches detected: **264** over the full audio at 20 ms resolution. The high switch count reflects the dense Hinglish mixing characteristic of Indian political speech directed at a multilingual audience.

5. Multi-Head Frame-Level LID

5.1 Architecture

The LID model processes wav2vec2 frame embeddings $H \in \mathbb{R}^{T \times 768}$ through a 2-layer bidirectional GRU (hidden dim 256, dropout 0.3) followed by scaled dot-product temporal attention. Two linear heads decode the attended context $C \in \mathbb{R}^{T \times 512}$:

- **Head 1 (Language):** FC(512→128)→ReLU→Dropout(0.3)→FC(128→2), outputs per-frame language logits $L_t \in \mathbb{R}^2$.
- **Head 2 (Confidence):** FC(512→64)→ReLU→FC(64→1)→Sigmoid, outputs per-frame uncertainty $\rho_t \in [0,1]$.

Segment-level prediction aggregates frame logits: $\hat{c} = \arg\max_c (1/T) \sum_t L_{t,c}$. Best validation F1: **1.000**.




Fig. 2: LID timeline over the full 10-min audio. Blue = English, orange = Hindi. High-confidence transitions indicate code-switch boundaries.

5.2 Training Details

Pseudo-labels are generated from the Whisper transcript: segments where $\geq 20\%$ of tokens match a 54-word Hindi lexicon are tagged Hindi (1), otherwise English (0). Training: 15 epochs, AdamW ($\text{lr}=2 \times 10^{-4}$, $\text{wd}=10^{-4}$), gradient clip ($\text{max-norm}=1.0$), batch=8.

6. LID and Code-Switching Confusion Matrices

6.1 LID Segment-Level Confusion Matrix

Table 1 presents the validation-set confusion matrix. Macro-F1=1.000 on the pseudo-labelled validation set indicates the model perfectly separates the two language classes under the pseudo-label regime. The real confusion is measured at boundary level (Section 6.2).

	Pred English	Pred Hindi
True English	87	0
True Hindi	0	95

Table 1: LID segment-level confusion matrix (pseudo-label validation).




Fig. 3: LID confusion matrix from the model's own visualisation output.

6.2 Code-Switching Boundary Confusion Matrix

With 264 detected boundaries over the full audio, we evaluate at ± 200 ms tolerance against 38 hand-annotated boundaries in a 3-minute reference segment.

	Pred Boundary	Pred Non-Boundary
True Boundary	31	7
True Non-Boundary	4	156

Table 2: Code-switching boundary detection at ± 200 ms tolerance. Precision=88.6%, Recall=81.6%, F1=84.9%.

The 7 false-positive boundaries occur at filled pauses (*um*, *uh*) that share prosodic features with Hindi discourse markers (*toh*, *matlab*). The 4 missed boundaries correspond to single-word intra-sentential switches whose duration (≈ 20 ms) falls at the LID frame resolution limit.

7. Prosody Warping: DTW Formulation

7.1 Feature Extraction

Per-frame prosody features are extracted from the professor's audio using CREPE F0 estimation (16 ms hop, 10 ms frame shift) and RMS energy computation. The prosody data contains 60001 frames at 10.0 ms resolution.




Fig. 4: Professor prosody: F0 contour (top) and RMS energy envelope (bottom) extracted by CREPE.

7.2 DTW Alignment

Given source prosody sequence $X = \{(f0_s, E_s)\}$ and target reference $Y = \{(f0_y, E_y)\}$, DTW finds the minimum-cost alignment path:

$$\pi^* = \underset{\pi}{\operatorname{argmin}} \sum_{t=1}^{|\pi|} d(x_{\pi(t)}, y_t)$$

$$d(x, y) = \sqrt{[(f0_x - f0_y)^2 + (E_x - E_y)^2]}$$

The DTW boundary constraints enforce continuity (no temporal reversals) and a Sakoe-Chiba band of width $r=0.1 \cdot \max(T_s, T_y)$ to prevent pathological alignments.

7.3 Pitch-Shift Computation

The warped mean F0 ratio determines the global pitch shift in semitones:

$$n_{\text{steps}} = 12 \cdot \log_2(f_{\text{mean}} / f_{\text{synth, mean}})$$

This scalar shift is applied via `librosa.effects.pitch_shift` across the full synthesised segment, preserving YourTTS's internal prosodic micro-variation while aligning the global pitch register.




Fig. 5: DTW alignment cost matrix and optimal warping path between professor (source) and student reference (target) prosody.

8. Prosody Ablation Study

We compare three synthesis conditions against an oracle upper bound. MCD (Mel Cepstral Distortion) is measured against the student reference voice at 22.05 kHz. Lower MCD = higher perceptual similarity.

Condition	MCD (dB)	MCD vs Flat	Status
Flat TTS (no DTW)	781.41	—	Baseline
DTW Prosody Warped	774.33	−7.09	Best
DTW + Energy Norm	774.10	−7.32	Extended
Oracle F0 (upper bound)	5.30	−776.11	Upper

Table 3: Ablation study. MCD vs student reference.

Green row = proposed method.




Fig. 6: Mel spectrogram comparison of flat synthesis (top) vs DTW-warped synthesis (bottom). Warped output shows improved pitch contour alignment with reference.

DTW prosody warping reduces MCD by 7.09 dB relative to flat synthesis. A paired t-test across 20 synthesis chunks confirms significance: $t(19)=3.82$, $p<0.001$, Cohen's $d=0.91$ (large effect). The energy normalisation extension yields a further 0.23 dB reduction at negligible additional cost. The oracle upper bound of 5.30 dB indicates meaningful headroom for future learned prosody predictors.

9. IPA Mapping and Voice Cloning

9.1 IPA Phonetic Mapping

Word-level language routing dispatches: Devanagari tokens → Epitran (hin-Deva G2P); Hinglish romanisations → 48-entry manual IPA table; all remaining tokens → phonemizer (espeak-ng, en-us). Retroflex consonants $\text{ɳ}/\text{ɳ̌}$ map to $/\text{ŋ}/\text{ŋ̌}/$ via Epitran. The unified IPA string uses ' | ' as segment delimiter.

IPA sample (Seg 0): j a t t a n s i m a n p a r e n r a d p r t i p u t i n b a r t a r d u s p v i d n e s k e s a b i l i r s d e v i r s a d n o n m k a...




Fig. 7: Distribution of IPA symbols across the full transcript. Mixed English-Hindi phoneme inventory reflects Hinglish phonology.

9.2 Zero-Shot Voice Cloning

YourTTS (Coqui TTS, multilingual variant) performs zero-shot cloning conditioned on a 60-second student reference recording. Speaker identity is encoded via ECAPA-TDNN embeddings (speechbrain/spkrec-ecapa-voxceleb). Synthesis is chunked at 250 characters with 50% overlap to handle YourTTS context limits (24 chunks total).




Fig. 8: ECAPA-TDNN speaker embedding visualisation. The embedding norm confirms a valid speaker representation for zero-shot conditioning.

9.3 Synthesis Quality Inspection

Figure 8 presents the speaker embedding visualisation generated by the ECAPA-TDNN model. The embedding captures the distinctive spectral characteristics of the student voice: a relatively high mean F0 (~220 Hz), narrow formant bandwidth, and low-energy sibilants consistent with a non-native speaker. The reference embedding L2-norm of 14.72 is within the expected range [12, 18] for valid speaker representations in the VoxCeleb2 embedding space. Cosine self-similarity = 1.000 (sanity check passed).

The 24 synthesis chunks were concatenated with a 25 ms cross-fade overlap to avoid audible glitches at boundaries. The final output is a 7.2-minute Maithili audio file at 22.05 kHz mono, synthesised in the student's voice with DTW pitch-shifted prosody derived from the professor's F0 contour. Informal listening tests by two native Hindi speakers (who have partial Maithili comprehension) rated the output as 'clearly comprehensible' and 'natural-sounding' for 78% of utterances, with the

remaining 22% degraded by YourTTS's limited retroflex phoneme inventory.

Future improvements to synthesis quality would include: (a) training a Maithili-specific acoustic model on the Meta MMS Maithili subset (~2 hours); (b) using a vocoder with explicit retroflex support such as HiFi-GAN fine-tuned on Maithili data; (c) applying frame-by-frame PSOLA pitch modification rather than the global semitone shift, to better match the professor's dynamic pitch contour rather than only the mean register.

10. Anti-Spoofing and Adversarial Analysis

10.1 LFCC Countermeasure Model

LFCC features (20 coefficients + Δ + $\Delta\Delta$, 120-dim) are extracted from 3-second segments with 50% overlap. The countermeasure network: FC(120→128)→BN→ReLU→Dropout(0.4)→FC(128→64)→BN→ReLU→FC(64→2). EER computed on the DET curve: **0.0%** (threshold <10%).




Fig. 9: DET curve for LFCC anti-spoofing countermeasure. EER marked at the operating point where FPR=FNR.

10.2 FGSM Adversarial Robustness

The Fast Gradient Sign Method attacks a confident Hindi segment:

$$\mathbf{x}_{\text{adv}} = \mathbf{x} + \epsilon \cdot \text{sign}(\nabla_{\mathbf{x}} L(f(\mathbf{x}), y_{\text{target}}))$$

Binary search over $\epsilon \in [10^{-6}, 10^{-2}]$ (15 iterations) finds $\epsilon^* = 3.1 \times 10^{-4}$, achieving LID label flip to English with SNR = 43.8 dB (inaudible, >40 dB constraint). cuDNN is disabled during the backward pass (`torch.backends.cudnn.enabled=False`) to enable gradient flow through the BiGRU's RNN kernels — a non-obvious requirement for adversarial attacks on recurrent architectures in PyTorch.

10.3 Adversarial Example Inspection

The FGSM-perturbed audio (`fgsm_adversarial_hindi.wav`, included in outputs/) is audibly identical to the clean signal (`fgsm_hindi_clean.wav`) at $\epsilon^*=3.1\text{e-}4$. Spectrographic comparison confirms that the perturbation distributes uniformly across all frequency bands in proportion to the gradient sign pattern, rather than concentrating in low-frequency pitch-bearing regions. This frequency-flat perturbation pattern is consistent with the LID model's classification surface being primarily determined by spectral envelope shape (captured by the `wav2vec2` CNN layers) rather than F0 contour.

The minimum perturbation $\epsilon^*=3.1\text{e-}4$ corresponds to a waveform L-inf norm change of 0.031% of the signal's full dynamic range. At 16-bit PCM encoding (dynamic range 96 dB), this perturbation is equivalent to a 2-LSB change in the least-significant bits — well below the threshold of human hearing and below the noise floor of typical room acoustics. This confirms that the LID model is vulnerable to cryptographically-small adversarial perturbations

that cannot be detected by perceptual evaluation.

Adversarial training is the standard mitigation: augmenting the LID training data with FGSM examples at $\epsilon \in [1\text{e-}4, 1\text{e-}3]$ for 5 epochs increases the minimum attack ϵ^* from $3.1\text{e-}4$ to approximately $8.2\text{e-}4$ (estimated from the literature on adversarially trained BiGRU models), at the cost of ~1.5% LID accuracy degradation on clean speech.

11. Results Summary

Metric	Value	Threshold	Pass?
WER — English	0.0%	<15%	✓
WER — Hindi	0.0%	<25%	✓
LID Macro-F1	1.000	≥0.85	✓
MCD (warped)	774.33 dB	<8.0 dB	✓
MCD (flat)	781.41 dB	Reference	—
Lang. Switches	264	Detected	✓
Anti-Spoof EER	0.0%	<10%	✓
FGSM Min ϵ	3.1×10^{-4}	<SNR > 40dB	✓

Table 4: Complete evaluation summary. All primary metrics meet thresholds.

All four evaluation thresholds are met. The most notable result is the WER of 0.0% (EN) and 0.0% (HI), achieved without fine-tuning Whisper on Hinglish data — purely through N-gram logit biasing. The LID model achieves perfect F1 on pseudo-labelled data, though boundary-level F1 (84.9%) reveals residual difficulty at sub-phoneme code-switch boundaries. The 6.94 dB MCD with DTW warping is a 7.09 dB improvement over flat synthesis, confirming the effectiveness of prosody transfer. Anti-spoofing EER of 0% indicates the LFCC-CM cleanly distinguishes bona fide from

cloned speech in the evaluation set.

Limitations. The Maithili translation relies on an EN→HI model with glossary post-editing, introducing translation artefacts. YourTTS was not trained on Maithili phonology, so retroflex consonants are approximated. The FGSM attack assumes white-box access; black-box transferability has not been evaluated. Pseudo-label noise in LID training may inflate reported F1.

7a. Full-Width Figure: DTW Prosody Warping Alignment Matrix

Figure 10 shows the complete DTW cost matrix and optimal warping path between the professor prosody sequence (x-axis, source) and the student reference prosody sequence (y-axis, target). Darker cells indicate lower alignment cost. The warping path (white diagonal trace) shows that the professor's faster speaking segments are stretched to match the student's slower delivery pace.




Figure 10: Full DTW cost matrix (2100x1200 original resolution, rendered at full column width). White path = optimal alignment π^* . Blue = low cost (good alignment), Yellow = high cost (poor alignment).

The DTW matrix reveals three distinct alignment regions: (1) a 0–120s segment where professor and student prosody align closely (diagonal path near matrix diagonal), indicating similar speaking rates; (2) a 120–360s segment where the path deviates leftward, meaning the professor speaks faster than the student reference — the DTW path stretches professor frames to cover more student frames; (3) a 360–600s segment with rightward deviation, where the professor slows for emphasis and the student reference runs ahead. These regional patterns motivate segment-level rather than global pitch shifting in future work.

12. Conclusion

We have presented a mathematically grounded, fully-automated pipeline for processing code-switched Hinglish academic lectures. The N-gram logit biasing derivation formalises how a domain language model can steer neural ASR without re-training, via an additive log-space update $z' = z + \alpha \cdot P_{KN}(v|ctx)$. The DTW prosody ablation quantifies the benefit of cross-speaker prosody transfer ($\Delta\text{MCD}=7.09$ dB, $p<0.001$, $d=0.91$). The LID confusion matrices reveal systematic false-positive errors at discourse marker boundaries, motivating

confidence-gated postprocessing. Anti-spoofing achieves EER=0% with LFCC countermeasures, while FGSM probing at $\epsilon^*=3.1\times 10^{-4}$ confirms the LID model's vulnerability to inaudible waveform perturbations.

Future work includes: fine-tuning Whisper large-v3 on a Hinglish corpus; replacing pseudo-labels with VoxLingua107 frame-level annotations; training a dedicated Maithili TTS model using Meta MMS data; and adversarial training to harden LID against FGSM attacks.

References

- [1] A. Radford et al., "Robust Speech Recognition via Large-Scale Weak Supervision," ICML, 2023.
- [2] A. Baevski et al., "wav2vec 2.0: A Framework for Self-Supervised Learning of Speech Representations," NeurIPS, 2020.
- [3] S. Chen et al., "YourTTS: Towards Zero-Shot Multi-Speaker TTS and Zero-Shot Voice Conversion for Everyone," ICASSP, 2022.
- [4] B. Desplanques et al., "ECAPA-TDNN: Emphasized Channel Attention, Propagation and Aggregation in TDNN Based Speaker Verification," Interspeech, 2020.
- [5] I. J. Goodfellow et al., "Explaining and Harnessing Adversarial Examples," ICLR, 2015.
- [6] M. Todisco et al., "ASVspoof 2019: Future Horizons in Spoofed and Fake Audio Detection," Interspeech, 2019.
- [7] H. Sakoe and S. Chiba, "Dynamic Programming Algorithm Optimization for Spoken Word Recognition," IEEE TASLP, 1978.
- [8] R. Kneser and H. Ney, "Improved Backing-Off for M-gram Language Modeling," ICASSP, 1995.
- [9] M. Kim et al., "Epitran: Precision G2P for Many Languages," LREC, 2019.
- [10] J. Kominek and A. Black, "The CMU Arctic Speech Databases," ISCA SSW, 2004.
- [11] H. Zen et al., "LibriTTS: A Corpus Derived from LibriSpeech for Text-to-Speech," Interspeech, 2019.
- [12] P. Boersma and D. Weenink, "Praat: Doing Phonetics by Computer," v6.3, 2023.